

# BLENDED\_LEARNING

---

## Implementation-of-Stochastic-Gradient-Descent-SGD-Regressor

---

### AIM:

---

To write a program to implement Stochastic Gradient Descent (SGD) Regressor for linear regression and evaluate its performance.

### Equipments Required:

---

1. Hardware – PCs
2. Anaconda – Python 3.7 Installation / Jupyter notebook

### Algorithm

---

- 1.Import Necessary Libraries: Import required libraries such as pandas, numpy, matplotlib, sklearn for the implementation.
- 2.Load the Dataset: Load the dataset (e.g., CarPrice\_Assignment.csv) using pandas.
- 3.Data Preprocessing: Drop unnecessary columns (e.g., 'CarName', 'car\_ID'). Handle categorical variables using `pd.get_dummies()`.
- 4.Split the Data: Split the dataset into features (X) and target variable (Y). Split the data into training and testing sets using `train_test_split()`.
- 5.Standardize the Data: Standardize the feature data (X) and target variable (Y) using `StandardScaler()` to ensure they have mean=0 and variance=1.
- 6.Create the SGD Regressor Model: Initialize the SGD Regressor model with `max_iter=1000` and `tol=1e-3`.
- 7.Train the Model: Fit the model to the training data using the `fit()` method.
- 8.Make Predictions: Use the trained model to predict the target values for the test set.
- 9.Evaluate the Model: Calculate performance metrics like Mean Squared Error (MSE) and R-squared score using `mean_squared_error()` and `r2_score()`.
- 10.Display Model Coefficients: Display the model's coefficients and intercept.
- 11.Visualize the Results: Create a scatter plot comparing actual vs predicted prices.

12.End: The program finishes by displaying the evaluation metrics, model coefficients, and a visual representation of the predictions.

## Program:

---

```
/*  
Program to implement SGD Regressor for linear regression.  
Developed by: GANESH PRABHU J  
RegisterNumber: 212223220023  
*/  
  
# Importing necessary libraries  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
from sklearn.model_selection import train_test_split  
from sklearn.linear_model import SGDRegressor  
from sklearn.metrics import mean_squared_error, r2_score  
from sklearn.preprocessing import StandardScaler  
  
# Load the dataset  
data = pd.read_csv('CarPrice_Assignment.csv')  
print(data.head())  
print("\n\n")  
print(data.info())  
  
# Data preprocessing  
# Dropping unnecessary columns and handling categorical variables  
data = data.drop(['CarName', 'car_ID'], axis=1)  
data = pd.get_dummies(data, drop_first=True)  
  
# Splitting the data into features and target variable  
x = data.drop('price', axis=1)  
y = data['price']  
  
# Standardizing the data  
scaler = StandardScaler()  
x = scaler.fit_transform(x)  
y = scaler.fit_transform(np.array(y).reshape(-1, 1)).ravel()  
  
# Splitting the dataset into training and testing sets  
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=  
  
# Creating the SGD Regressor model  
sgd_model = SGDRegressor(max_iter=1000, tol=1e-3)  
  
# Fitting the model on the training data  
sgd_model.fit(x_train, y_train)
```

```
# Making predictions
y_pred = sgd_model.predict(x_test)

# Evaluating model performance
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

# Print evaluation metrics
print()
print()
print("Mean Squared Error:", mse)
print("R-squared Score:", r2)

# Print model coefficients
print("\n\n")
print("Model Coefficients")
print("Coefficients:", sgd_model.coef_)
print("Intercept:", sgd_model.intercept_)

# Visualizing actual vs predicted prices
print("\n\n")
plt.scatter(y_test, y_pred)
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices using SGD Regressor")
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red') # Perfect line
plt.show()
```

## Output:

---

 [simple linear regression model for predicting the marks scored](https://github.com/ganeshprabhu2005/BLENDED_LEARNING_Implementation-of-Stochastic-Gradient-Descent-SGD-Regressor/tree/main)

## LOAD THE DATASET:

```

car_ID  symboling      CarName  fueltype  aspiration  doornumber  \
0      1          3    alfa-romero giulia      gas      std      two
1      2          3    alfa-romero stelvio    gas      std      two
2      3          1  alfa-romero Quadrifoglio  gas      std      two
3      4          2      audi 100 ls        gas      std      four
4      5          2      audi 100ls        gas      std      four

carbody  drivewheel  enginelocation  wheelbase  ...  enginesize  \
0  convertible    rwd      front      88.6  ...    130
1  convertible    rwd      front      88.6  ...    130
2   hatchback    rwd      front      94.5  ...    152
3      sedan     fwd      front      99.8  ...    109
4      sedan     4wd      front      99.4  ...    136

fuelsystem  boreratio  stroke  compressionratio  horsepower  peakrpm  citympg  \
0      mpfi      3.47    2.68          9.0      111    5000    21
1      mpfi      3.47    2.68          9.0      111    5000    21
2      mpfi      2.68    3.47          9.0      154    5000    19
3      mpfi      3.19    3.40         10.0      102    5500    24
4      mpfi      3.19    3.40          8.0      115    5500    18

highwaympg  price
0          27  13495.0
1          27  16500.0
2          26  16500.0
3          30  13950.0
4          22  17450.0

[5 rows x 26 columns]

```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                205 non-null    int64
1   symboling              205 non-null    int64
2   CarName               205 non-null    object
3   fueltype              205 non-null    object
4   aspiration             205 non-null    object
5   doornumber            205 non-null    object
6   carbody               205 non-null    object
7   drivewheel            205 non-null    object
8   enginelocation        205 non-null    object
9   wheelbase             205 non-null    float64
10  carlength             205 non-null    float64
11  carwidth              205 non-null    float64
12  carheight             205 non-null    float64
13  curbweight            205 non-null    int64
14  enginetype            205 non-null    object
15  cylindernumber        205 non-null    object
16  enginesize            205 non-null    int64
17  fuelsystem            205 non-null    object
18  boreratio             205 non-null    float64
19  stroke                205 non-null    float64
20  compressionratio      205 non-null    float64
21  horsepower            205 non-null    int64
22  peakrpm               205 non-null    int64
23  citympg               205 non-null    int64
24  highwaympg            205 non-null    int64
25  price                 205 non-null    float64
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB
None
```



## EVALUATION METRICS:

Mean Squared Error: 0.1342093761275513

R-squared Score: 0.8588105274203248

## MODEL COEFFICIENTS:

Model Coefficients

Coefficients: [ 5.95324665e-02 1.04621666e-01 -1.50496499e-02 2.02200858e-01

1.59477431e-02 1.70910766e-01 3.28878621e-01 -4.72514838e-02

-8.08308399e-02 -1.31203362e-02 1.31879946e-01 2.72801451e-04

-7.62997727e-03 -2.10543720e-02 -1.98904900e-02 7.10134411e-03

7.59905720e-03 -4.16349472e-02 -9.69939527e-02 -1.84480888e-02

-6.89098925e-02 -5.18959127e-02 7.79779092e-02 2.22160463e-01

9.60416043e-03 -8.76711936e-02 8.80296430e-02 7.85793361e-05

7.17936131e-03 3.48416168e-03 -4.29482616e-02 -1.59998058e-01


-9.00523043e-02 1.65784211e-03 -3.40666404e-02 3.48416168e-03

-1.81138301e-02 3.05131038e-03 1.98904900e-02 1.65784211e-03

-3.87934514e-02 -3.60520608e-02 -1.76071141e-02]

Intercept: [-0.02368564]

## VISUALIZATION OF ACTUAL VS PREDICTED VALUES:



## Result:

---

Thus, the implementation of Stochastic Gradient Descent (SGD) Regressor for linear regression has been successfully demonstrated and verified using Python programming.